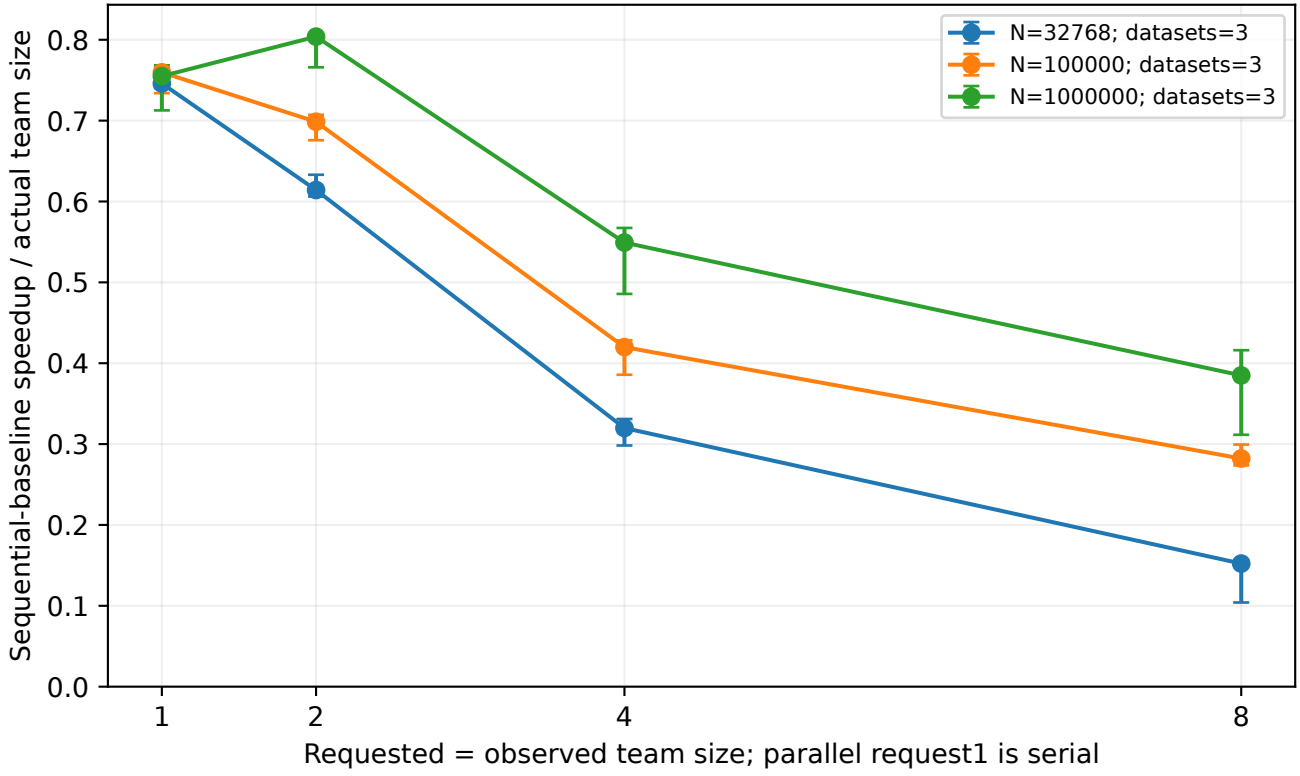
nearly_sorted: median of dataset summaries; min–max spread
N=32768; datasets=3
N=100000; datasets=3
N=1000000; datasets=3
Sequential-baseline speedup / actual team size
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
1
2
4
8
Requested = observed team size; parallel request1 is serial